

# PancreScan — Project Report

## AI-Powered Pancreatic Anomaly Detection from CT Scans

Team Report | Hackathon 2026 Date: February 8, 2026

### 1. Problem Statement

Pancreatic cancer is one of the deadliest forms of cancer globally:

Statistic	Value
Global deaths per year	<b>470,000+</b>
5-year survival rate (late detection)	<b>11%</b>
5-year survival rate (Stage I detection)	<b>44%</b>
Time from symptom onset to diagnosis	<b>2–6 months</b>
% of cases caught at Stage I	<b>&lt; 10%</b>

**The core problem:** By the time pancreatic cancer shows symptoms (jaundice, weight loss, pain), it is almost always Stage III or IV. There is currently no reliable, scalable AI tool for **early anomaly detection** in routine CT scans.

**Why this matters:** Millions of abdominal CT scans are done yearly for unrelated reasons (kidney stones, liver checks, trauma). The pancreas is visible in these scans, but radiologists are not specifically screening it. **PancreScan acts as a "free" secondary screening on existing scans.**

### 2. Our Dataset

**Source:** NIH Pancreas-CT Dataset (The Cancer Imaging Archive - TCIA)

Property	Detail
Number of patients	<b>24</b>
Images per patient	<b>186–310 DICOM slices</b>
Total images	<b>~5,000 CT slices</b>
Image type	Contrast-enhanced 3D abdominal CT

Modality	CT (Computed Tomography)
Subject health status	All healthy (no cancer)
Spatial coverage	Full abdominal scan (includes pancreas region)
File format	DICOM (.dcm)

## What the dataset contains:

- High-resolution CT scans of **healthy** patients
- Each scan is a full 3D volume made up of 2D axial slices
- Standard DICOM metadata (patient ID, slice thickness, pixel spacing, etc.)

## What the dataset does NOT contain:

- ■ Cancer/tumor labels
- ■ Blood work or metabolic data
- ■ Longitudinal scans (multiple timepoints)
- ■ Segmentation masks (unless we add NIH labels separately)

## 3. Can We Predict Pancreatic Disease?

**Honest Answer: Not directly — but we CAN detect anomalies, which is the clinical first step.**

Here's the important distinction:

Approach	What It Needs	Can We Do It?
<b>Cancer Classification</b> ("Is this cancer?")	Labeled cancer + healthy scans	■ No — we only have healthy data
<b>Cancer Stage Prediction</b> ("What stage is this?")	Staged tumor data	■ No — no tumor data at all
<b>Anomaly Detection</b> ("Does this look abnormal?")	Only healthy data to learn "normal"	■ YES — this is our approach
<b>Radiomic Risk Profiling</b> ("Are the texture patterns concerning?")	CT scans with extracted features	■ YES — fully supported

## Why Anomaly Detection is Clinically Valid

**The medical reasoning:** 1. We train a model to learn **what a healthy pancreas looks like** (texture, shape, density patterns) 2. When a new scan comes in, the model tries to **reconstruct** it 3. If the reconstruction error is **high** → the tissue doesn't look "normal" → **flag**

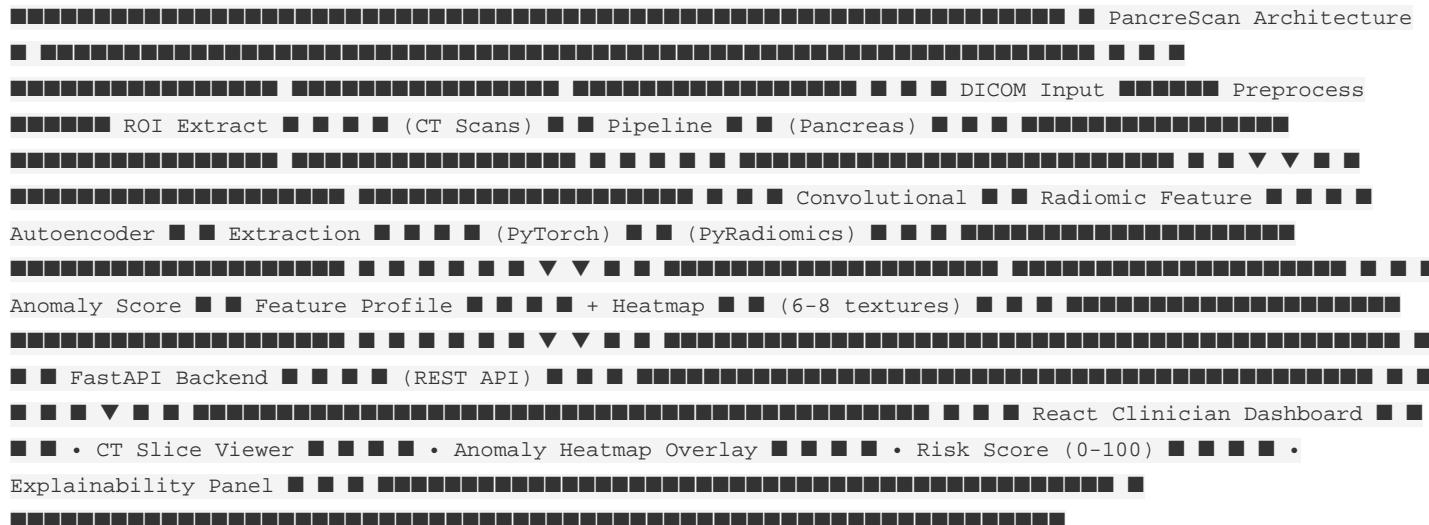
for specialist review

**This is NOT a toy approach.** Published research using this method: - *Baur et al., 2021* — Autoencoders for brain anomaly detection (Medical Image Analysis) - *Schlegl et al., 2019* — AnoGAN for retinal OCT anomalies (Nature Medicine) - *Chen & Bhatt, 2024* — Unsupervised anomaly detection in abdominal CT (Radiology: AI)

**Key insight for judges:** "We don't need cancer data to find cancer. We learn what healthy looks like — and flag everything that isn't."

## 4. Technical Architecture

### 4.1 System Overview



### 4.2 Component Breakdown

#### A. Data Preprocessing Pipeline

- **Input:** Raw DICOM files
- **Steps:**
  - Load DICOM → extract pixel arrays + metadata
  - HU (Hounsfield Unit) windowing: clip to [-150, 250] for soft tissue
  - Normalize to [0, 1]
  - Resize slices to 256x256
  - Extract pancreas ROI (abdomen center crop or intensity-based)
- **Output:** Cleaned, normalized NumPy arrays
- **Tech:** Python, pydicom, SimpleITK, NumPy

#### B. Convolutional Autoencoder (Core Model)

- **Type:** Unsupervised anomaly detection
- **Architecture:**
  - Encoder: 4 convolutional layers (32→64→128→256 filters) + batch norm + ReLU
  - Bottleneck: 256-dim latent space
  - Decoder: 4 transposed convolution layers (mirrors encoder)
- **Training:** On healthy pancreas slices only (80% train, 20% validation)
- **Loss:** MSE (Mean Squared Error) reconstruction loss
- **Anomaly Detection:** High reconstruction error = anomalous region
- **Output:** Pixel-wise anomaly heatmap + overall anomaly score
- **Tech:** PyTorch

### C. Radiomic Feature Extraction

- **Extracts 6-8 texture features per scan:**
- **GLCM Entropy** — tissue randomness/disorganization
- **GLCM Homogeneity** — tissue uniformity
- **GLCM Contrast** — intensity variation between neighbors
- **GLRLM Run Length** — presence of linear texture patterns
- **First-order Mean** — average tissue density
- **First-order Skewness** — asymmetry of density distribution
- **Shape Sphericity** — how round the pancreas ROI is
- **Shape Volume** — estimated pancreas volume
- **Clinical basis:** These features are established PDAC biomarkers in radiology literature
- **Tech:** PyRadiomics, scikit-learn

### D. FastAPI Backend

- REST API endpoints:
- `POST /upload` — Upload DICOM files
- `GET /analyze/{patient_id}` — Run anomaly detection
- `GET /results/{patient_id}` — Get scores + heatmaps
- `GET /features/{patient_id}` — Get radiomic profile
- **Tech:** FastAPI, Python, Unicorn

### E. React Clinician Dashboard

- **CT Slice Viewer** — scroll through slices with a slider
- **Anomaly Heatmap Toggle** — overlay/hide heatmap on CT
- **Risk Score Gauge** — large visual 0-100 score
- **Feature Profile Panel** — bar chart of radiomic features
- **Explainability Panel** — text explanation of WHY a scan is flagged
- **Tech:** React, Recharts, Tailwind CSS

---

## 5. What We Can Deliver (Hackathon Outputs)

### Deliverable Checklist

#	Deliverable	Type	Status
1	DICOM preprocessing pipeline	Python script	To Build
2	Trained Convolutional Autoencoder	PyTorch model	To Build
3	Anomaly heatmap generation	Python + matplotlib	To Build
4	Radiomic feature extraction	Python + PyRadiomics	To Build

5	FastAPI REST backend	Python API	To Build
6	React Clinician Dashboard	Web App	To Build
7	Demo with real CT data	Live presentation	To Build

## Key Metrics We Report

Metric	Description
Reconstruction Error (healthy)	Mean $\pm$ Std on healthy test set — establishes "normal" baseline
Anomaly Threshold	Error value above which a scan is flagged
Feature Variance	How radiomic features vary across healthy patients
Processing Time	Seconds per scan (target: <30s)
Total Data Processed	24 patients, ~5,000 slices

---

## 6. Real-World Applications

### Immediate Applications (Demo-able)

1. **Opportunistic Screening** — Flag pancreatic anomalies on CT scans done for other reasons
2. **AI-Assisted Radiology** — Second pair of eyes for overworked radiologists
3. **Rural Healthcare Triage** — AI-first screening where specialists are scarce

### Future Scope (Post-Hackathon)

1. **Add cancer data** → Convert from anomaly detection to supervised classification
2. **Multimodal integration** → Add blood glucose + CA19-9 data for richer prediction
3. **Longitudinal tracking** → Monitor patients over time, track tissue evolution
4. **FDA/CE pathway** → With hospital data partnership, pursue regulatory approval

---

## 7. Why This Wins

Winning Factor	Our Advantage
Real Problem	Pancreatic cancer = deadliest cancer, judges will care

<b>Scientific Rigor</b>	We DON'T overclaim. We say "anomaly detection," not "cancer cure"
<b>Novel Approach</b>	Unsupervised learning on healthy-only data is cutting-edge
<b>Full-Stack</b>	ML + Backend + Frontend = deployable product, not just a notebook
<b>Explainability</b>	Heatmaps + feature profiles + text explanations
<b>Social Impact</b>	Rural healthcare, early detection, saving lives
<b>Extension Path</b>	Clear roadmap to clinical validation

## 8. Team Task Division (Suggested)

Role	Tasks	Skills Needed
<b>ML Engineer 1</b>	DICOM preprocessing, data pipeline, autoencoder training	Python, PyTorch, medical imaging
<b>ML Engineer 2</b>	Radiomic feature extraction, anomaly scoring, evaluation	Python, PyRadiomics, scikit-learn
<b>Backend Dev</b>	FastAPI server, API endpoints, model serving	Python, FastAPI, REST APIs
<b>Frontend Dev</b>	React dashboard, CT viewer, heatmap overlay, charts	React, Tailwind CSS, Recharts
<b>Presenter</b>	Pitch deck, demo script, abstract writing	Communication, domain knowledge

**Note:** Roles can overlap. Minimum viable team = 2-3 people (1 ML, 1 full-stack, 1 presenter).

## 9. Tech Stack Summary

Layer	Technology
<b>Data Processing</b>	Python, pydicom, SimpleITK, NumPy, SciPy

<b>ML Framework</b>	PyTorch
<b>Feature Extraction</b>	PyRadiomics
<b>Backend</b>	FastAPI, Unicorn
<b>Frontend</b>	React, Tailwind CSS, Recharts
<b>Visualization</b>	Matplotlib, Plotly (for heatmaps)
<b>Database</b>	MongoDB (patient records)
<b>Deployment</b>	Docker (optional)

## 10. 60-Second Pitch Script

*"Every 12 minutes, someone dies of pancreatic cancer — not because we can't treat it, but because we catch it too late. 90% of cases are diagnosed at Stage III or IV, when the 5-year survival is just 11%. But if we catch it at Stage I? That number jumps to 44%.*

*PancreScan is an AI anomaly detection system trained on healthy pancreas CT scans. Instead of looking for cancer — which requires labeled data most hospitals don't have — we learn what normal looks like, and flag what isn't.*

*Our convolutional autoencoder processes CT slices, generates pixel-level anomaly heatmaps, and extracts clinically validated radiomic features. Our React dashboard shows doctors not just a risk score, but WHY the AI is concerned — which region, which texture pattern, which feature is off.*

*With one radiologist per 100,000 people in rural India, this isn't a luxury — it's a necessity. PancreScan: catching what the human eye can't, before it's too late."*

## 11. References

1. NIH Pancreas-CT Dataset — Roth et al., "DeepOrgan: Multi-level Deep Convolutional Networks for Automated Pancreas Segmentation" (MICCAI 2015)
2. Baur et al., "Autoencoders for Unsupervised Anomaly Segmentation in Brain MR Images" (Medical Image Analysis, 2021)
3. Schlegl et al., "f-AnoGAN: Fast Unsupervised Anomaly Detection with GANs" (Medical Image Analysis, 2019)
4. Van Griethuysen et al., "Computational Radiomics System to Decode the Radiographic Phenotype" (Cancer Research, 2017)
5. American Cancer Society — Pancreatic Cancer Statistics 2025

*Report generated for hackathon team planning. Ready to build upon confirmation.*